



CLOSET

Three-dimensional articulated figures for Typst



Version 0.1.0

Contents

1. Overview	3
1.1. Installation	3
1.2. First figure	3
2. Three-dimensional model	3
2.1. Constructing a pose manually	4
2.2. Composing gesture modifiers	4
2.3. Built-in poses	5
2.4. Controlled variation	5
3. Hand gestures	7
3.1. Canonical gestures	7
3.2. Beckoning from three directions	7
3.3. Electromagnetic right-hand rule	8
3.4. Volumetric tube rendering	8
3.5. Defining a hand pose	9
3.6. Angles, landmarks, and variation	9
4. Camera and projection	10
4.1. One pose, three views	10
4.2. Choosing a useful viewpoint	11
4.3. Framing and focal length	11
5. Calligraphic rendering	11
6. Imported OpenSim motion	12
7. API reference	12
7.1. hand, hand-pose, and hand-gestures	12
7.2. hand-angles, hand-joints, and vary-hand	12
7.3. pose	13
7.4. Gesture modifiers	13
7.5. human	13
7.6. vary-pose	13
7.7. camera	13
7.8. project	13
7.9. joints	13
7.10. pose-from-degrees and constrain-pose	13
8. Troubleshooting	13

1. Overview

Closet draws compact human figures from an articulated model. A pose contains three-dimensional bone directions. The package evaluates those directions into joint positions, projects the joints through a perspective camera, and sends the resulting paths to Premetadated's calligraphic pen engine.

The pipeline is:

```
3D pose directions
-> 3D forward kinematics
-> perspective camera
-> 2D projected paths
-> calligraphic nib envelopes
-> CeTZ vector output
```

This separation is important. Changing the camera does not alter the body or its pose; it only changes how the same 3D joints appear on the page.

1.1. Installation

For local package use, install both repositories under Typst's local package directory:

```
{data-dir}/typst/packages/local/closet/0.1.0
{data-dir}/typst/packages/local/premetadated/0.1.0
```


On macOS, {data-dir} is ~/Library/Application Support. Import the package as:

```
#import "@local/closet:0.1.0": human
```

From this repository, examples can instead use `#import "lib.typ"`.

1.2. First figure

```
#import "@local/closet:0.1.0": human
#human(pose: "yoga-warrior-two", scale: 1.2)
```



*Warrior two with the default
camera and pen*

`human` returns a CeTZ canvas, so it can be placed in normal Typst grids, figures, boxes, and alignment containers.

2. Three-dimensional model

Closet uses right-handed (x, y, z) coordinates:

- x runs horizontally across a front-facing body;
- y carries depth, toward or away from the viewer;
- z points upward.


A canonical pose stores `torso` and `side` directions plus one direction for each upper and lower arm and leg. An optional `head` direction lets the head move independently; when omitted, it follows the torso.

Directions need not be unit length because the kinematics normalizes them before applying fixed anatomical segment lengths.


2.1. Constructing a pose manually

Use `pose(...)` to construct a complete pose. Every argument has a standing default, so specify only the directions that differ. This example makes a twisting crouch with one arm reaching forward and the other moving back:

```
#let custom = closet.pose(  
  torso: (0.0, -0.32, 0.95),  
  head: (0.15, -0.2, 0.97),  
  left-upper-arm: (0.35, -0.85, 0.4),  
  left-lower-arm: (0.1, -1.0, 0.1),  
  right-upper-arm: (-0.45, 0.75, -0.5),  
  right-lower-arm: (-0.15, 0.9, -0.4),  
  left-upper-leg: (0.3, -0.7, -0.65),  
  left-lower-leg: (-0.1, 0.25, -0.97),  
  right-upper-leg: (-0.3, 0.5, -0.82),  
  right-lower-leg: (0.1, -0.2, -0.98),  
)  
  
#closet.human(pose: custom)
```



Front: eye (0, -8, 2.4)



Side: eye (8, 0, 2.4)

Each arm and leg direction is local to its segment: the lower-arm vector starts at the elbow, and the lower-leg vector starts at the knee. Positive z points upward. Test a new pose from several cameras because depth components can be hidden in a single projection.

`joints(pose)` evaluates a pose without drawing it. It returns `pelvis`, `neck`, and `head` points plus three-point chains for each arm and leg. Every returned point is a 3D tuple.

2.2. Composing gesture modifiers

Gesture helpers accept either a built-in name or a pose dictionary, and return a new pose dictionary. They can therefore be nested without changing the original. Here a standing figure takes a step, leans, extends the right arm, waves, and tilts its head:

```
#let greeting = closet.tilt-head(  
  closet.wave-arms(  
    closet.stretch-arms(  
      closet.bend-torso(  
        closet.walk-legs("standing", amount: 0.45),  
        amount: 0.12,  
      ),  
      amount: 0.5,  
      side: "right",  
    ),  
    amount: 0.85,  
    side: "right",  
  ),  
)
```


```

    amount: -0.18,
  )
  #closet.human(pose: greeting)


```



Standing



Walk legs



Bend + stretch



Wave + tilt


raise-arms, stretch-arms, spread-legs, and wave-arms use an amount from 0 to 1. tilt-head, bend-torso, and walk-legs accept signed values from -1 to 1; changing the sign reverses the tilt, bend, or stride. Arm helpers accept side: "left", "right", or "both". These helpers blend 3D directions and preserve segment lengths, but they are posing conveniences, not collision detection or physiological validation.

2.3. Built-in poses


The package includes standing, sitting, cross-legged sitting, thinking, walking, running, downward dog, child's pose, warrior one, warrior two, tree pose, squatting, push-up, pull-up, walking on hands, jumping on one foot, and jumping with open legs. Their stable names are available as keys in `poses`; see `assets/previews/poses.png` for the visual index.



Thinking




One-foot jump




Warrior two



Tree



Push-up



Hand walk

2.4. Controlled variation

`vary-pose(pose, variation:, seed:)` adds small deterministic offsets to the canonical 3D bone directions. It accepts either a built-in pose name or a pose dictionary and returns a new dictionary suitable for `human`, `joints`, or additional editing.

```

#let people = range(5).map(seed =>
  vary-pose("walking", variation: 0.09, seed: seed)
)

#for person in people {
  human(pose: person, camera: profile-camera)
}

```

The same operation is available directly on `human`:

```

#human(pose: "walking", variation: 0.09, seed: 3)

```



Seed 0



Seed 1



Seed 2



Seed 3



Seed 4

The integer seed makes output reproducible. Equal inputs produce identical figures, while different seeds redistribute each limb slightly. Variation is bounded to 0.04 to 0.35 ; 0 returns the canonical pose, and 0.04 to 0.12 usually gives useful scene diversity without changing the gesture. Torso and body-side axes receive a smaller perturbation than limbs. All directions are normalized after perturbation, so bone lengths remain fixed.

3. Hand gestures

The hand model follows the standard 21-landmark topology: one wrist, four thumb landmarks, and four landmarks for each finger. It is a small procedural model, not a general anatomical simulator. Canonical gestures store finger curl and spread, thumb opposition, global spread, and wrist orientation. Closet expands those controls into constrained joint angles and fixed-length 3D bones, then uses the same camera projection and Premetadated pen as human.

3.1. Canonical gestures




Open palm



Fist



Pointing



V / peace




Thumbs-up



Pinch / OK



Three fingers



Rock



Beckoning

```
#hand(  
  gesture: "peace",  
  handedness: "right",  
  camera: camera(  
    eye: (2.4, -7.0, 3.0),  
    target: (0.0, 0.0, 0.65),  
    focal-length: 6.5,  
  ),  
)
```

The built-in names are open-palm, fist, pointing, peace, thumbs-up, pinch, three-fingers, rock, beckoning, and right-hand-rule; ok aliases pinch. Handedness is anatomical, not screen-relative. With the palm facing the viewer, a right hand's thumb appears on the viewer's right; from the back it appears on the left. Set handedness: "left" to reflect the complete 3D skeleton before wrist rotation.

3.2. Beckoning from three directions

The beckoning gesture curls the index finger partway while the remaining fingers stay more tightly flexed. These are three projections of the same 3D gesture. Its finger shading is reduced to one short hatch per phalanx so the joint bends remain clear.



Front



Three-quarter




Side

```
#hand(
  gesture: "beckoning",
  finger-hatch-count: 1,
  camera: side-view,
)
```

3.3. Electromagnetic right-hand rule

right-hand-rule forms a right-handed three-finger triad. The extended index finger is y , the naturally bent middle finger is z , and the thumb is x . The three arrows follow the corresponding 3D finger directions rather than their apparent directions on the page. Set `axes: true` to superpose labeled arrow vectors using the same camera projection.



Right hand: thumb x , index y , middle z

```
#hand(
  gesture: "right-hand-rule",
  axes: true,
  finger-hatch-count: 1,
  camera: view,
)
```

3.4. Volumetric tube rendering

The default `render: "tube"` sends each 3D finger chain and thicker metacarpal paths through Premetadated's tube renderer, then adds an oblate ellipsoid for the palm. The solids currently overlap with visible seams; a future Larnt-level union can remove those boundaries. Use `render: "skeleton"` to inspect the projected landmark centerlines. Sparse light-directed hatching distinguishes phalanges and depth without stippling; set `shading: false` for outlines only. Both modes use the same gesture, landmarks, handedness, and camera.



Projected skeleton



Volumetric tubes

```
#hand(
  gesture: "peace",
  render: "tube",
  tube-radius: 0.05,
  tube-sides: 10,
  shading: true,
  camera: view,
)
```

`tube-radius` is measured in the hand model's 3D units. `tube-sides` controls cross-section smoothness and must be at least six. `shading` enables sparse line hatching on the tubes and palm; it never uses stippling. `finger-hatch-count` controls marks per phalanx and defaults to 2; use 1 for sparse shading or 0 for unshaded fingers. The tube renderer derives its field of view from Closet's camera focal length, so changing `eye`, `target`, `up`, or `focal-length` works consistently in both modes.

3.5. Defining a hand pose


`hand-pose(...)` supplies straight-finger defaults. Finger `curl` runs from 0 to 1; per-finger `spread` runs from -1 to 1; global `spread` runs from 0 to 1.5. Thumb `opposition` runs from 0 to 1. Wrist pitch and yaw are limited to $\pm 45^\circ$, and roll to $\pm 90^\circ$.

```
#let custom = hand-pose(
  thumb: (curl: 0.35, spread: 0.1, opposition: 0.8),
  index: (curl: 0.05, spread: -0.35),
  middle: (curl: 0.2, spread: 0.15),
  ring: (curl: 0.75, spread: 0.1),
  pinky: (curl: 0.9, spread: 0.2),
  spread: 0.7,
  wrist: (pitch: 8deg, yaw: -12deg, roll: 5deg),
)

#hand(gesture: custom)
```



Right hand



Mirrored left hand

3.6. Angles, landmarks, and variation


`hand-angles(gesture, variation:, seed:)` exposes the expanded MCP/PIP/DIP angles. The four long fingers use maximum flexions of 70° , 95° , and 65° ; thumb CMC/MCP/IP maxima are 40° , 55° , and 60° . Because all three joints derive from one varied curl value, their motion remains correlated.

`hand-joints(gesture, variation:, seed:, handedness:)` returns the 21 named 3D points: `wrist`; `thumb-cmc`, `thumb-mcp`, `thumb-ip`, `thumb-tip`; and `MCP`, `PIP`, `DIP`, and `tip` points for index, middle, ring, and pinky. `hand-topology` exposes the five connected chains. This matches MediaPipe's topology and terminology, without adding MediaPipe or an ML runtime dependency.

```
#let landmarks = hand-joints("pointing")
#let variants = range(5).map(seed =>
  vary-hand("pointing", variation: 0.07, seed: seed)
)
```



Seed 0



Seed 1



Seed 2



Seed 3



Seed 4

Variation is deterministic and bounded to $0..0.25$. It perturbs each finger's shared curl, lateral spread, thumb opposition, and wrist orientation, then clamps every value before angle expansion. Segment lengths never change.

4. Camera and projection

Create a perspective camera with:


```
#let view = camera(
  eye: (3.8, -8.0, 2.8),
  target: (0.0, 0.0, 0.2),
  up: (0.0, 0.0, 1.0),
  focal-length: 7.0,
)
```

`eye` is the camera position. `target` is the world-space point at the center of the image. `up` establishes image vertical. `focal-length` controls perspective magnification: increasing it enlarges the projected figure when camera distance is unchanged.

The implementation constructs a look-at basis from these values. For point p , eye e , forward vector f , camera-right vector r , and camera-up vector u , depth is $d = (p - e) \cdot f$. Projection is proportional to $((p - e) \cdot \frac{r}{d}, (p - e) \cdot \frac{u}{d})$. Points must remain in front of the camera.

4.1. One pose, three views

The following three figures all use exactly `poses.sitting-cross-legged`. Only the camera position changes.



Front



Three-quarter



Side

The front view emphasizes bilateral symmetry. The three-quarter view reveals which knees and hands advance in depth. The side view compresses the body's width but exposes front-to-back placement. This is why a pose that looks ambiguous from one direction can become immediately legible from another.

```
#let pose = poses.sitting-cross-legged

#human(pose: pose, camera: camera(eye: (0, -8, 2.4)))
#human(pose: pose, camera: camera(eye: (5, -7, 3.0)))
#human(pose: pose, camera: camera(eye: (8, 0, 2.4)))
```

4.2. Choosing a useful viewpoint

Use a front view for poses whose defining feature is left-right symmetry, such as warrior two, tree pose, squatting, and jumping. Use a profile view for poses whose silhouette depends on forward motion or floor contact, such as sitting, walking, running, downward dog, child's pose, and push-up. A three-quarter view is useful when depth relationships matter but both sides must remain visible.

The supplied `front-camera`, `profile-camera`, and `three-quarter-camera` are convenient defaults, not semantic constraints. A pose may face another world direction, and a custom `camera` can always be more informative.

4.3. Framing and focal length


Move `eye` and `target` to frame the subject. Keep `target` near the torso for upright poses and lower it for floor poses. If the figure is too small, increase `focal-length`, increase `human(scale:)`, or move the eye closer. Changing focal length alters perspective and magnification; changing `scale` only changes the final drawing size and pen dimensions.



Focal length 4.5



Focal length 7



Focal length 10


5. Calligraphic rendering

Every bone and the head outline are rendered by `premetadated.stroke.nib-stroke`. The default tangent-following nib has a broad axis of `0.050`, a fine axis of `0.015`, and an offset of `24deg`. Nib dimensions scale with the figure.


Pass any `Premetadated` pen dictionary to override it:

```
#human(
  pose: "yoga-warrior-two",
  pen: (
    mode: "calligraphic",
    offset: 12deg,
    samples: ((arclength: 0, a: 0.055, b: 0.009)),
  ),
  stroke: rgb("352b25"),
)
```


Larger separation between `a` and `b` creates stronger contrast between broad and fine directions. The nib offset rotates that contrast relative to each path's tangent.



Fine



Broad



Rotated

6. Imported OpenSim motion

`tools/import_opensim.py` is an offline adapter. It downloads pinned Gait2354 model and inverse-kinematics data, then writes compact model ranges and paired gait frames to `data/gait2354.json`.

```
python3 tools/import_opensim.py
typst compile --root . examples/opensim.typ build/opensim.pdf
```

`pose-from-degrees(base, values)` converts imported sagittal hip and knee angles into 3D leg directions. Use a canonical 3D pose as the upper-body base:

```
#let data = json("data/gait2354.json")
#let sample = data.samples.first()
#let pose = pose-from-degrees(poses.walking, sample.pose_degrees)
#human(pose: pose, limits: data.limits, camera: profile-camera)
```



right leg forward at 0.6833333 s



left leg forward at 1.3 s

The synchronized hip and knee values belong to one observed time step. Joint ranges alone do not guarantee a realistic combined pose, and the Gait2354 ranges are model domains rather than universal physiological safety limits.

7. API reference

7.1. hand, hand-pose, and hand-gestures

`hand(gesture:, variation:, seed:, handedness:, render:, tube-radius:, tube-sides:, shading:, finger-hatch-count:, axes:, scale:, stroke:, pen:, camera:)` renders a procedural hand. `axes: true` overlays the coordinate triad for right-hand-rule. `render` accepts "skeleton" or "tube". `hand-pose(...)` creates compact finger controls, and `hand-gestures` contains the canonical dictionaries.

7.2. hand-angles, hand-joints, and vary-hand

`hand-angles(gesture, variation:, seed:)` expands controls into constrained angles. `hand-joints(gesture, variation:, seed:, handedness:)` evaluates the 21 3D landmarks. `vary-hand(gesture, variation:, seed:)` returns a reproducibly perturbed control dictionary. `hand-topology` and `hand-joint-limits` expose the connectivity and limits used by the model.

7.3. pose

`pose(torso:, side:, head:, left-upper-arm:, left-lower-arm:, right-upper-arm:, right-lower-arm:, left-upper-leg:, left-lower-leg:, right-upper-leg:, right-lower-leg:)` constructs a complete pose dictionary. All arguments are optional and default to standing. `head: none` follows the torso direction.

7.4. Gesture modifiers

`tilt-head`, `bend-torso`, `raise-arms`, `stretch-arms`, `spread-legs`, `walk-legs`, and `wave-arms` each return a modified copy of a named or custom pose. Functions that affect arms accept `side:` where appropriate. Their results can be passed to another modifier, `human`, `joints`, or `vary-pose`.

7.5. human

`human(pose:, variation:, seed:, scale:, stroke:, pen:, limits:, camera:)` renders one figure. `pose` accepts a built-in name or complete 3D pose dictionary. `stroke` is the ink color. `pen` accepts a `Premetadated` pen; `none` selects the package default. `limits` optionally clamps imported OpenSim leg angles. `camera` controls the projection.

7.6. vary-pose

`vary-pose(pose, variation:, seed:)` returns a reproducibly perturbed 3D pose. The accepted variation range is 0..0.35 and the seed must be an integer.

7.7. camera

`camera(eye:, target:, up:, focal-length:)` validates and returns a camera dictionary. The package exports `default-camera`, `front-camera`, `profile-camera`, and `three-quarter-camera`.

7.8. project

`project(point, camera:)` maps one 3D point to a 2D pair. It is useful when labels or other graphics must align with a figure joint.

7.9. joints

`joints(pose)` evaluates forward kinematics and returns the 3D body landmarks without rendering.

7.10. pose-from-degrees and constrain-pose

These helpers bridge compact OpenSim-derived data to the 3D renderer. They are primarily intended for generated sagittal gait samples.

8. Troubleshooting

- **Package not found:** install both `closet:0.1.0` and `premetadated:0.1.0` in the same local package namespace.
- **A pose looks flat:** move the camera away from a principal axis, for example from (0, -8, 2.4) to (4, -7, 3).
- **A limb is hidden:** rotate the camera around `target`; do not modify the pose merely to compensate for one viewpoint.
- **The figure is clipped:** move `eye` farther away or reduce `scale`.
- **Projection reports a point behind the camera:** move `eye` away from the model and keep `target` between the camera and figure.
- **Calligraphy is too subtle:** increase nib axis `a` or reduce `b`.